MAT377 Lecture Notes

Arky!! :3c

'24 Fall Semester

Contents

§1 Day 1: Introduction to Probability (Sep. 4, 2024)

Link to textbook.

We start with a sampler problem that on the surface, seems unrelated to probability. Let $v_1, \ldots, v_n \in \mathbb{R}^n$ be unit vectors on the unit sphere, i.e.t $\|v_i\| = 1$. If we are to pick $\varepsilon_i = \{-1, 1\}$ at random, what is our expectation on how large will

$$\sum_{i=1}^{n} \varepsilon_i v_i$$

be? We could brute force and average out over all probabilities as follows,

$$\frac{1}{2^n} \sum_{\varepsilon \in \{-1,1\}^n} \left(\sum_{i=1}^{n} \varepsilon_i v_i \right) = \sum_{i=1}^{n} \left(\frac{1}{2^n} \underbrace{\sum_{\varepsilon \in \{-1,1\}^n} \varepsilon_i}_{=0} \right) v_i = 0.$$

Now, consider that

$$\begin{aligned} \frac{1}{2^n} \sum_{\varepsilon \in \{-1,1\}^n} \left| \sum_{i=1}^{n} \varepsilon_i v_i \right|^2 &= \frac{1}{2^n} \sum_{\varepsilon \in \{-1,1\}^n} \sum_{i_1, i_2 = 1}^{n} \varepsilon_{i_1} \varepsilon_{i_2} \langle v_{i_1}, v_{i_2} \rangle \\ &= \sum_{i_1, i_2 = 1}^{n} \left(\frac{1}{2^n} \sum_{\varepsilon \in \{-1,1\}^n} \varepsilon_{i_1} \varepsilon_{i_2} \right) \langle v_{i_1}, v_{i_2} \rangle \end{aligned}$$

To simplify the bracketed summation, we could consider the following two cases:

- If $i_1 \neq i_2$, we would have that

$$\frac{1}{2^n} \sum_{\varepsilon \in \{-1,1\}^n} \varepsilon_{i_1} \varepsilon_{i_2} = \frac{2^{n-2}}{2^n} \sum_{\substack{\varepsilon_{i_1} \in \{-1,1\} \\ \varepsilon_{i_2} \in \{-1,1\}}} \varepsilon_{i_1} \varepsilon_{i_2} = 0.$$

- If $i_1 = i_2$, we would have

$$\frac{1}{2^n} \sum_{\varepsilon \in \{-1,1\}^n} \varepsilon_{i_1} \varepsilon_{i_2} = \frac{2^{n-1}}{2^n} \sum_{\varepsilon_i \in \{-1,1\}} \varepsilon_{i_1} \varepsilon_{i_2} = 1.$$

By linearity of expectation, we obtain

$$\frac{1}{2^n} \sum_{\varepsilon \in \{-1,1\}^n} \left| \sum_{i=1}^{n} \varepsilon_i v_i \right|^2 = n,$$

and

$$\frac{1}{2^n} \sum_{\varepsilon \in \{-1,1\}^n} \left| \sum_{i=1}^{n} \varepsilon_i v_i \right| = c\sqrt{n},$$

where c is a positive real constant.

We now abstract a few probability terms;

- Ω is a sample space, i.e. the set of possible outcomes.
- Let P denote probability, i.e. a mapping of subsets of Ω to $[0,1]$ (read: probability of getting these subsets of Ω); the probability of an event ε out of S occuring is given by $P(\varepsilon \in S, S \subset \{-1,1\}^n) = \frac{1}{|S|}$, assuming that each event in S is equally likely. With this, we have three important properties of P to define:
  1. $P(\Omega) = 1$; the chance of an event in the probability space happening is 1.
  2. Let $\mathcal{F}$ be a collection of subsets $A_1, \ldots, A_n$. Then
  $$P\left(\bigcup_{i=1}^{n} A_i\right) = \sum_{i=1}^{n} P(A_i),$$
  given that $A_i \cap A_j = \emptyset$ for all $1 \leq i, j \leq n$. This is linearity of expectation.
  3. $P(A^C) = 1 - P(A)$, which is a property of set complement.
- When our collection $\mathcal{F}$ of subsets of Ω satisfy the following properties, we call it a σ-algebra:
  1. $\emptyset \in \mathcal{F}$,
  2. Closed under countable union: $A_n \in \mathcal{F} \implies \bigcup_{i=1}^{n} A_i \in \mathcal{F}$,
  3. Closed under complement: $A_i \in \mathcal{F} \implies A_i^C \in \mathcal{F}$.

  In a finite sample space, the power set $\mathcal{F} = \mathcal{P}(\Omega)$ is one such example of a σ-algebra. As an example, let $\Omega = [0,1)$; then $P([a,b]) = b - a$ (wlog, let $a < b$). Now, let $\mathcal{F} = \mathcal{P}([0,1))$. Define the equivalence $x \sim y$ if $x - y \in \mathbb{Q}$.
- We now introduce the axiom of choice; Let A be a set containing one element of each equivalence class from the above defined equivalence. Consider $\tau_q A := \{A\} + q$; let us claim that
  $$\bigcup_{q \in \mathbb{Q}} \tau_q A = [0,1),$$
  which is a countable union of $[0,1)$, since $\mathbb{Q}$ is countable. We have that $P(A) = P(\tau_q A)$ because intervals don't change size under shifting by q. However, observe that
  $$P([0,1)) = P\left(\bigcup_{q \in \mathbb{Q}} \tau_q A\right) = \sum_{q \in \mathbb{Q}} P(\tau_q A).$$
  Then either
  $$\begin{aligned} P(A) = 0 &\implies P(\tau_q A) = 0 \implies P([0,1)) = 0, \text{or} \\ P(A) \neq 0 &\implies P(\tau_q A) \to \infty \implies P([0,1)) \to \infty, \end{aligned}$$
  which doesn't make sense (for now). This shows that we need to pick our σ-algebra properly; observing that the intersection of two σ-algebras is also a σ-algebra, it is appropriate to let $\mathcal{F}$ be the smallest σ-algebra containing $[a,b)$. This is called a *Borel Set*.[1]

[1] this is confusing. ill check later

§2 Day 2: Expectations and Distributions (Sep. 9, 2024)

Course administrative details first; starting next week, office hours will be held on Monday from 11:15am to 12:15pm. Recap of last lecture:

- A probability space Ω is the set of all possible outcomes of an "experiment," i.e. a countable set of individual events $\{\omega_1, \dots, \omega_n\}$ (we will cover continuous probability later on).
- $\mathcal{F} = \mathcal{P}(\Omega)$ is the set of all subsets of Ω.
- $P(A) = \sum_{\omega \in \Omega} P(\omega)$ is the probability of an outcome in $A \in \mathcal{F}$ occurring.

A random variable X is a function $\Omega \to \mathbb{R}$, aka the measurement of the event, and the expectation of the random variable, EX, is given by $\sum_{\omega \in \Omega} X(\omega)P(\omega)$. Is expectation well behaved? No. For example, consider the St. Petersburg Paradox; suppose you are playing a game in the casino; every time you flip a coin, your prize money doubles if it lands on heads (read: double or nothing lfg!!!). Then we may consider the set of outcomes to be the number of consecutive heads, i.e.

$$\begin{aligned} \Omega &= \{1, 2, 3, \dots\}, \\ P(n) &= \frac{1}{2^n}, \\ X(n) &= 2^n. \end{aligned}$$

Clearly, the chance of getting n heads in a row is 2^{-n}, and assuming your prize money started at 1 dollar, you would win 2^n dollars for said n heads. Taking the expectation of this game, we find

$$EX = \sum_{n=1}^{\infty} 2^n \cdot \frac{1}{2^n} = \sum 1 = \infty.$$

It doesn't make sense to expect to win infinite amounts of money from this game[2] unless you had unlimited wealth to start with. With this in mind, we insist on

$$\sum_{\omega \in \Omega} |X(\omega)| \, P(\omega) < \infty$$

within the context of this class.

Theorem 2.1 (Linearity of Expectation)**.** X is linear; i.e., $E[ax + by] = aEx + bEy$.

We start with a lemma:

Lemma 2.2. Let us have a bijective map $\pi : \mathbb{N} \times \mathbb{N} \to \mathbb{N}$. Then

$$\sum_{n=1}^{\infty} c_n = \sum_{n=1}^{\infty} \sum_{m=1}^{\infty} C_{\pi(n,m)}$$

if all $c_n \geq 0$ or if either side is absolutely convergent.

[2] martingale strat lfg.,..

We proceed to prove this with casework.

- Suppose $c_n \geq 0$; then using the bijective nature of π, we may choose large enough N, M such that
$$\sum_{n=1}^{K} c_n \leq \sum_{n=1}^{N}\sum_{m=1}^{M} c_{\pi(n,m)}$$
for any choice of K. Conversely, we may pick
$$\sum_{n=1}^{N}\sum_{m=1}^{M} c_{\pi(n,m)} \leq \sum_{n=1}^{K} c_n$$
for any N, M by picking $k \geq \max_{\substack{1\leq n\leq N\\ 1\leq m\leq M}}\{\pi(n,m)\}$. Now, let $M \to \infty$; we have
$$\sum_{n=1}^{N}\sum_{m=1}^{\infty} c_{\pi(n,m)} \leq \sum_{n=1}^{\infty} c_n,$$
then let $N \to \infty$ to get
$$\sum_{n=1}^{\infty}\sum_{m=1}^{\infty} c_{\pi(n,m)} \leq \sum_{n=1}^{\infty} c_n.$$
As per earlier, we also see that LHS is greater or equal to RHS, which implies equality. □

- Now, suppose $\sum_{n=1}^{\infty} |c_n| < \infty$. Let $c_n = a_n - b_n$, where $a_n = c_n 1(c_n \geq 0)$ and $b_n = c_n 1(c_n < 0)$. Then we obtain
$$\sum_{n=1}^{\infty} a_n = \sum_{n=1}^{\infty}\sum_{m=1}^{\infty} a_\pi(n,m), \quad \sum_{n=1}^{\infty} b_n = \sum_{n=1}^{\infty}\sum_{m=1}^{\infty} b_\pi(n,m)$$
as per our proof above. Summing both, we conclude that equality holds for absolute convergence as well. □

For now, let X take values $\{a_1, a_2, \dots\}$ (countably many). Consider
$$P'(a_n) = P(X = a_n) = P(\underbrace{\{\omega \mid X(\omega) = a_n\}}_{X^{-1}(a_n)})$$
as the probably of a pre-image (or, $P' = P \circ X$). We see that P' is a probability on $\mathbb{R}$ (concentrated on $\{a_1, a_2, \dots\}$), and $0 \leq P'(a_n) \leq 1$ for any n; from now, we will call P' a *distribution* of X. Here are some examples of distributions:

- The Bernoulli distribution: let $0 \leq p \leq 1$. Then consider a coin with p chance to land on heads, and $1-p$ on tails; then Ber_p is given by $\Omega = \{H, T\}$,
$$X(H) = 1, P(X = 1) = p,$$
$$X(T) = 0, P(X = 0) = 1 - p.$$

- Flip N coins, with $X = \{0, 1, \dots, N\}$ being the number of heads we obtain. Then
$$P(X = \ell) = \binom{N}{\ell} p^\ell (1-p)^{N-\ell},$$

and the expected value is given by

$$EX = \sum_{\ell=0}^{N} \ell \binom{N}{\ell} p^{\ell}(1-p)^{N-\ell}.$$

Using linearity of expectation, we see $EX = EX_1 + \cdots + EX_n = Np$ by separating each coinflip.

Expectation enjoys the change of variables property;[3]

$$EX = \sum_{\omega \in \Omega} X(\omega)P(\omega) = \sum_{n=1}^{\infty} a_n P'(a_n).$$

To see this, consider partioning the probability space Ω into $X^{-1} = \{\omega_{nm} \mid 1 \leq m \leq M_n\}$ in terms of their measurement from X (where $X(\omega_{ni}) = X(\omega_{nj}) = a_n$ for any $1 \leq i, j \leq M_n$)[4], and write

$$\sum_{\omega \in \Omega} X(\omega)P(\omega) = \sum_{n=1}^{\infty} \sum_{m=1}^{M_n} X(\omega_{nm})P(\omega_{nm}), \qquad \text{(by Lemma)}$$

where we may note that mapping each individual $\omega \in \Omega$ to some index nm is bijective since it is a partition. We continue by writing

$$\begin{aligned} &= \sum_{n=1}^{\infty} \sum_{m=1}^{M_n} a_n P(\omega_{nm}) = \sum_{n=1}^{\infty} a_n \left(\sum_{m=1}^{M_n} P(\omega_{nm}) \right) \\ &= \sum_{n=1}^{\infty} a_n P(X = a_n), \end{aligned}$$

where we may note $P(X = a_n) = P'(a_n)$. □

The probability distribution of any given random variable X also approaches 0 at its tail. Specifically, we have that $\lim_{t \to \infty} P(x \geq t) = 0$. To prove this, we start by observing that $P(x \geq t)$ is monotone decreasing; consider

$$P(X \geq n) = \sum_{m=n}^{\infty} P(m \leq X < m+1).$$

Clearly, the sum is convergent, as the sum of probabilities is equal to 1. Using the fact that the tail of a convergent series approaches 0, we conclude that $P(m \leq X < m+1) \to 0$ as $m \to \infty$, and so $P(X \geq n) \to 0$ as $n \to \infty$.

[3] read: sum of value of outcome multiplied by the chance it occurs over all ω is the same as going over each value individually and multiplying the chance you roll into it

[4] read M_n as a counter of how many outcomes in Ω have the same measurement of a_n

Lemma 2.3 (Expectation of Random Variable in terms of Integral)**.** The expectation of a random variable X may be expressed as $EX = \int_0^\infty P(X \geq t)\,dt$ for $X \geq 0$.[5]

Let us start by considering the case where X takes integer values only;

$$\begin{aligned} EX = \sum_{n=1}^{\infty} nP(X=n) &= \sum_{n=1}^{\infty}\sum_{m=1}^{n} P(X=n) \\ &= \sum_{m=1}^{\infty}\sum_{n=m}^{\infty} P(X=n) \\ &= \sum_{m=1}^{\infty} P(X \geq m). \end{aligned}$$

For the general case, let us start by writing $a_n = \int_0^\infty 1(t \leq a_n)\,dt$ by the layer cake decomposition. Then

$$\begin{aligned} EX = \sum_{n=1}^{\infty} a_n P(X=a_n) &= \sum_{n=1}^{\infty}\left(\int_0^\infty 1(t\leq a_n)\,dt\right)P(X=a_n) \\ &\overset{(*)}{=} \int_0^\infty \left(\sum_{n=1}^{\infty} 1(t\leq a_n)P(X=a_n)\right)dt \qquad \text{(Fubini)} \\ &= \int_0^\infty \sum_{a_n \geq t} P(X=a_n)\,dt \\ &= \int_0^\infty P(X\geq t)\,dt \end{aligned}$$

To resolve $(*)$ without the use of Fubini's theorem, we may write

$$\begin{aligned} \sum_{n=1}^{\infty}\int_0^\infty 1(t\leq a_n)P(X<a_n)\,dt &= \sum_{n=1}^{\infty}\sum_{m=1}^{\infty}\int_{m-1}^{m} 1(t\leq a_n)P(X=a_n)\,dt \\ &= \sum_{m=1}^{\infty}\sum_{n=1}^{\infty}\int_{m-1}^{m} 1(t\leq a_n)P(X\geq a_n)\,dt \\ &= \sum_{m=1}^{\infty}\lim_{N\to\infty}\sum_{n=1}^{N}\int_{m-1}^{m} 1(t\leq a_n)P(X\geq a_n)\,dt \\ &= \sum_{m=1}^{\infty}\lim_{N\to\infty}\int_{m-1}^{m}\left(\sum_{n=1}^{N} 1(t\leq a_n)P(X\geq a_n)\right)dt \\ &= \sum_{m=1}^{\infty}\int_{m-1}^{m}\left(\lim_{N\to\infty}\sum_{n=1}^{N} 1(t\leq a_n)P(X\geq a_n)\right)dt \\ &= \int_0^\infty\left(\sum_{n=1}^{\infty} 1(t\leq a_n)P(X=a_n)\right)dt, \end{aligned}$$

which we conclude by removing the auxiliary summations, since it is enough to know that $\sum_{n=1}^{N} 1(t\leq a_n)P(X=a_n) \to \sum_{n=1}^{\infty} 1(t\leq a_n)P(X=a_n)$ uniformly in $t \in [m-1, m]$; i.e.,

$$\left|\sum_{n=N+1}^{\infty} 1(t\leq a_n)P(X=a_n)\right| \leq \sum_{n=N+1}^{\infty} P(X=a_n) \to 0$$

as $N \to \infty$ as per earlier (since the tail goes to 0).

[5] intuition: layer cake formula, but compile them together in level sets.

We also briefly went over examples multinomial distributions at the end of class;

- Suppose $X_1, \ldots, X_n$ are independent, and let $P(X_i = j) = p_j$ for $j = 1, \ldots, k$. Let $\Omega = \{n_1, \ldots, n_k), n_j \geq 0, n_1 + \cdots + n_k = n\}$ (read: k-sided dice rolled n times, where n_j denotes the number of times j came up). Then

$$P((n_1, \ldots, n_k)) = \binom{n}{n_1, \ldots, n_k} p_1^{n_1} p_2^{n_2} \ldots p_k^{n_k}.$$

- The geometric distribution; let $0 < p < 1$ denote the probability of getting a head, and let us toss a coin until we get a heads. Let the outcome of X denote the number of tosses it took. Then

$$P(X = n) = (1 - p)^{n-1} p,$$

  and we may check $\sum_{n=1}^{\infty} P(x = n) = 1$ by geometric series.

- The Poisson distribution; let $\lambda > 0$. Then

$$P(X = n) = \frac{\lambda^n}{n!} e^{-\lambda}$$

  for $n = 0, 1, 2, \ldots$.

§3 Day 3: Distributions, Stability Property, Moments (Sep. 11, 2024)

Recall the Poisson distribution

$$P(X=n)=\frac{\lambda^n}{n!}e^{-\lambda}$$

where $n=0,1,\dots$ and $\lambda>0$. We introduce the stability property: let us consider the independent random variables X_1, X_2 distributed as follows,

$$\begin{aligned}X_1 &\sim \text{Poiss}_{\lambda_1},\\ X_2 &\sim \text{Poiss}_{\lambda_2}.\end{aligned}$$

Then we have that $X_1+X_2\sim\text{Poiss}_{\lambda_1+\lambda_2}$. To prove this, write

$$\begin{aligned}P(X_1+X_2=n)&=\sum_{m=0}^{n}P(X_1=m,X_2=n-m)\\ &=\sum_{m=0}^{n}P(X_1=m)P(X_2=n-m)\\ &=\sum_{m=0}^{n}\frac{\lambda_1^m}{m!}e^{-\lambda_1}\frac{\lambda_2^{m-2}}{(m-n)!}e^{-\lambda_2}\\ &=\frac{1}{n!}\underbrace{\sum_{m=0}^{n}\frac{n!}{m!(n-m)!}\lambda_1^m\lambda_2^{n-m}}_{(\lambda_1+\lambda_2)^n}e^{-(\lambda_1+\lambda_2)}\\ &=\frac{1}{n!}(\lambda_1+\lambda_2)^n e^{-(\lambda_1+\lambda_2)}.\end{aligned}$$

Binomials also have a related property; let

$$\begin{aligned}X_1&\sim\text{Bin}(n_1,p),\\ X_2&\sim\text{Bin}(n_2,p).\end{aligned}$$

Given that X_1, X_2 are independent, we know that $X_1+X_2\sim\text{Bin}(n_1+n_2,p)$. To prove this, we may just write

$$\begin{aligned}X_1&=y_1+\cdots+y_{n_1},\\ X_2&=y_{n_1+1}+\cdots+y_{n_1+n_2},\\ X_1+X_2&=y_1+\cdots+y_{n_1+n_2}\sim\text{Bin}(n_1+n_2,p).\end{aligned}$$

Moreover, we also have $\text{Bin}(n,\frac{\lambda}{n})\overset{n\to\infty}{\to}\text{Poiss}_\lambda$. This is called the *law of little numbers*. To prove this, we have

$$\begin{aligned}\binom{n}{k}\left(\frac{\lambda}{n}\right)^k\left(1-\frac{\lambda}{n}\right)^{n-k}&=\frac{\lambda^k}{k!}\underbrace{\frac{n(n-1)\dots(n-k+1)}{n^k}}_{\to 1}\underbrace{\left(1-\frac{\lambda}{n}\right)^{n-k}}_{\to e^{-\lambda}}\\ &\overset{n\to\infty}{\to}\frac{\lambda^k}{k!}e^{-\lambda}=\text{Poiss}_\lambda.\end{aligned}$$

There are two examples of Poisson distributions that we will go over: shark attacks and radioactive decay. (but we didn't go over it ig?)

Theorem 3.1 (Doeblin)**.** Let X_i be independent random variables distributed by Ber_{p_i}, where $0 < p_i < 1$. Let us have $S_n = X_1 + \cdots + X_n$ with $\lambda = p_1 + \cdots + p_n$. Then

$$\left| P(S_n \in A) - \sum_{n \in A} \frac{\lambda^n}{n!} e^{-\lambda} \right| \leq \sum_{i=1}^{n} p_i^2,$$

where $A \subset \{0, 1, \dots\}$.

To prove this, let y be a random variable where $y \sim \text{Poiss}_p$, then $P(y = 0) = e^{-p} > 1 - p$. Define $\Omega^\perp = \{-1, 0, 1, 2, \dots\}$; then we have $P_p(-1) = 1 - p, P_p(0) = e^{-p} - 1 + p, P_p(k) = \frac{p^k}{k!} e^{-p}$ for $k = 1, 2, 3, \dots$. Moreover, define

$$X(\omega) = \begin{cases} 0 & \omega = -1 \\ 1 & \omega \geq 0 \end{cases}, \quad y(\omega) = \begin{cases} 0 & \omega = 0, 1 \\ \omega & \omega \geq 1 \end{cases}.$$

Then $P(x = y) = 1 - p + pe^{-p} \geq 1 - p + p(1 - p) = 1 - p^2$, so $P(x = y) \leq p^2$. Take $\Omega = (\Omega_+)^n$ and $X_i(\omega) = X(\omega)$, and let us have

$$P(\omega) = \prod_{i=1}^{n} P_{p_i}(\omega_i)$$

where X_i are independently distributed by Ber_p, and y_i are independently distributed by Poiss_{p_i}. Finally, let us have

$$S_n = X_1 + \cdots + X_n,$$
$$S_n' = y_1 + \cdots + y_n,$$

then $P(S_n \neq S_n') \leq \sum_{i=1}^{n} P(X_i \neq y_i) \leq \sum_{i=1}^{n} p_i^2$, which means $S_n' \sim \text{Poiss}_\lambda$.[6] □

We now define *moments*. For a random variable X, $E(X) = \sum_{\omega \in \Omega} X(\omega) P(\omega)$, as long as $E|X| < \infty$. Then moments are given by EX^n where $n = 1, 2, \dots$. (??) For example, let $X = \text{Poiss}_\lambda$. Then we have for $n = 1$,

$$\begin{aligned} EX &= \sum_{n=0}^{\infty} n \frac{\lambda^n}{n!} e^{-\lambda} \\ &= \sum_{n=0}^{\infty} \frac{\lambda^n}{(n-1)!} e^{-\lambda} \\ &= \sum_{n=0}^{\infty} \frac{\lambda^{n+1}}{n!} e^{-\lambda} \\ &= \lambda. \end{aligned}$$

For $n = 2$, we have

$$\begin{aligned} EX^2 &= \sum_{n=0}^{\infty} n^2 \frac{\lambda^n}{n!} e^{-\lambda} \\ &= E(X(X-1)) + EX \\ &= \lambda^2 + \lambda. \end{aligned}$$

[6] reminder: review this proof, i'm stupid and don't really get it.

§4 Day 4: Independence and Dependence (Sep. 16, 2024)

Let (Ω, P) be our probability space. We define

$$P(A \mid B) := \frac{P(A \cap B)}{P(B)}$$

to be the conditional probability, i.e. probability of A given B, as long as $P(B) > 0$ (this is called Bayes' Rule). If $P(A \mid B) = P(A)$, then A is said to be independent of B. In particular, if $A_1, \ldots, A_n$ are independent, then

$$P\left(\bigcap_{i=1}^{n} A_i\right) = \prod_{i=1}^{n} P(A_i).$$

If the above is true for only pairs of events A_i, A_j, then we say that they are pairwise independent.

Let Ω_i, P_i be probability spaces, and consider $\Omega = \Omega_1 \times \cdots \times \Omega_n = \prod_{i=1}^{n} \Omega_i$, where we define a probability event in $\omega \in \Omega$ to be $(\omega_1, \ldots, \omega_n) = \omega$ with $\omega_i \in \Omega_i$. Specifically, we have

$$P(\omega) := \prod_{i=1}^{n} P_i(\omega_i).$$

For example, let $A = A_1 \times \cdots \times A_n$, and $A_i \in \Omega_i$. then

$$P(A) := \sum_{\omega \in A} P(\omega) = \sum_{\substack{\omega_i \in A_i \\ i=1,\ldots,n}} \prod_{i=1}^{n} P_i(\omega_i) = \prod_{i=1}^{n} \sum_{\omega_i \in A_i} P_i(\omega_i) = \prod_{i=1}^{n} P_i(A_i).$$

Let us have random variables $X_i : \Omega_i \to \mathbb{R}$ where $1 \le i \le n$. Then $X_i^{(\omega)} = f_i(\omega_i)$ are independent if $P(X_1 \in A_1, \ldots, X_n \in A_n) = \prod_{i=1}^{n} P(X_i \in A_i)$. In other words, $X_i^{-1}(A_i)$ are independent.[7] We may continue simplifying the expression as follows,

$$\prod_{i=1}^{n} P(X_i \in A_i) = \prod_{i=1}^{n} P_i(f_i(\omega_i) = x_i) = \prod_{i=1}^{n} P(X_i = x_i).$$

Now, suppose X, Y are independent and f, g are functions. Then we claim that $f(X), g(Y)$ are independent. To check this, let us write

$$\begin{aligned} P(f(X) = a, g(Y) = b) &= P(X = f^{-1}(a), y = g^{-1}(b)) \\ &= P(X \in f^{-1}(a)) P(Y \in g^{-1}(b)) \\ &= P(f(X) = a) P(g(Y) = b). \end{aligned}$$

We can also do this with grouping; let $\{1, \ldots, n\} = U_{k=1}^{m} I_k$ with I_k disjoint; i.e., we're sorting $[n]$ into disjoint subsets I_k. Then let $y_k = f_k(\{x_i\}_{i \in I_k})$ for some function $f_k : \mathbb{R}^{|I_k|} \to \mathbb{R}$, and we have that y_k are independent. To prove this, observe that

$$\begin{aligned} P(y_1 \in A_1, \ldots, y_m \in A_m) &= P(f_1 \in A_1, \ldots, f_m \in A_m) \\ &= P(\{X_i\}_{i \in I_1} \in f_1^{-1}(A_1), \ldots, \{X_i\}_{i \in I_m} \in f_m^{-1}(A_m)) \\ &= \prod_{j=1}^{m} P(\{X_i\}_{i \in I_j} \in f_j^{-1}(A_j)) \\ &= \prod_{j=1}^{m} P(y_j \in A_j). \end{aligned}$$

[7] note on board: for any $A_1, \ldots, A_n$ borel sets, intervals are enough, like $(-\infty, x_i]$. confusion?

We need to show that $P(\{X_i\}_{i\in I_1} = b_1, \{X_i\}_{i\in I_2} = b_2) = P(\{X_i\}_{i\in I_1} = b_1)P(\{X_i\}_{i\in I_2} = b_2)$; but as per earlier, this is true.

Now, suppose our random variables X_is are independently binomial distributed. Then

$$X_1 + \cdots + X_{m_1} \sim \text{Bin}(m_1, p) \sim \text{Poiss}_{\lambda_1}, \qquad (\lambda_1 = pm_1)$$

$$X_{m_1+1} + \cdots + X_{m_1+m_2} \sim \text{Bin}_{m_2,p} \sim \text{Poiss}_{\lambda_2} \qquad (\lambda_2 = pm_2)$$

We may combine the groupings above to get $X_1 + \ldots X_{m_1+m_2} \sim \text{Bin}(m_1 + m_2, p) \sim \text{Poiss}_{\lambda_1+\lambda_2}$.

Lemma 4.1. If X and Y are indepnendent and $E[|X|] < \infty$, $E[|Y|] < \infty$, then $E[XY] = E[X]E[Y]$.[8]

First, assume $X, Y > 0$. Let us directly write

$$\begin{aligned} E[XY] &= \sum_{\omega\in\Omega} X(\omega)Y(\omega)P(\omega) \\ &= \sum_{n,m} a_n b_m P(X = a_n, Y = b_m) \\ &= \sum_{n,m} a_n b_m P(X = a_n)P(Y = b_m) \\ &= \sum_{n} a_n P(X = a_n) \sum_{m} b_m P(Y = b_m) \\ &= E[X]E[Y]. \end{aligned}$$

In the case that the random variables are not necessarily non-negative, we may simple consider

$$\begin{aligned} X &= X1(X \geq 0) - |X|\,1(X < 0) = X_+ - X_-, \\ Y &= Y1(Y \geq 0) - |Y|\,1(Y < 0) = Y_+ - Y_-. \end{aligned}$$

However, do note that in the OPPOSITE direction that $E[XY] = E[X]E[Y]$ does NOT imply that X, Y are independent. It is true that $E[f(X)g(Y)] = E[f(X)]E[g(y)]$ for "lots of" f, g would imply that X, Y independent (if this is true for *all* f, g, then it is independent), but this is unreliable.

Using Fubini's theorem, we may consider X, Y on non-discrete probability spaces, and write

$$\begin{aligned} E[f(X,Y)] &= \sum_{n,m} f(a_n, b_m)P(X = a_n, Y = b_m) \\ &\overset{\text{if indep.}}{=} f(a_n, b_m)P(X = a_n)P(Y = b_m) \\ &\overset{\text{if "nice"}}{=} \sum_{n} \left[\sum_{m} f(a_n, b_m)P(Y = b_m) \right] P(X = a_n). \end{aligned}$$

We say that the above is "nice" if $f \geq 0$, or $E[f(x, y)] < \infty$, or

$$\sum_{n} \left[\sum_{m} |f(a_n, b_m)|\, P(Y = b_m) \right] P(X = a_n) < \infty.$$

[8] i'm sick of the no bracket nonsense

Alternatively, if we dont have our "nice" cases, we have

$$\sum_n \left[\sum_m f(a_n, b_m) P(Y = b_m \mid X = a_n) \right] P(X = a_n).$$

Now, we inntroduce the conditional distribution $P(y = b_m \mid X = a_n)$, where the distribution is Y given $X = a_n$. We can write the expectation

$$E[g(Y) \mid X = a_n] = \sum_m g(b_m) P(y = b_m \mid x = a_n),$$

i.e. the conditional expectation of $g(Y)$ given $X = a_n$. For example, let $X_1, X_2, \dots$ be i.i.d. Ber_p and N indep. Poiss_λ. Then $Y = X_1 + \cdots + X_N$ has

$$P(Y = k) = \sum_{n=0}^{\infty} P(Y = k, N = n) = \sum_{n=0}^{\infty} P(y = k \mid N = k) P(N = n).$$

If $N = n$, then $Y = X_1 + \cdots + X_n$, and we have

$$P(Y = k \mid N = n) = P(X_1 + \cdots + X_n = k) = \binom{n}{k} p^k (1-p)^{n-k}.$$

Specifically,

$$\begin{aligned} P(y) &= \sum_{n=0}^{\infty} \frac{n!}{k!(n-k)!} p^k (1-p)^{n-k} \frac{\lambda^n}{n!} e^{-\lambda} \\ &= \frac{(\lambda k)^k}{k!} \left(\sum_{n=0}^{\infty} \frac{(1-p)^{n-k}}{(n-k)!} \lambda^{n-k} \right) e^{-\lambda} \\ &= \frac{(\lambda p)^k}{k!} e^{-p\lambda} \sim \text{Poiss}_{p\lambda}. \end{aligned}$$

Now for another example; let $X_1, X_2, \dots$ be i.i.d. Ber_{y_2}; i.e. let $x_i \in \{0,1\}^{\mathbb{N}}$; let $x \in [0,1) = \Omega$, $X = 0, X_1, X_2, \dots$. Let P be on $[0,1)$. Then $P([a,b)) = b - a$ where $b > a$; we claim that they are i.i.d. Ber_{y_2}, which is proven by subdividing the intervals (whatever this means).

Let $P(X_1 = x_1, \dots, X_n = x_n)$. Then this is equal to

$$\begin{aligned} &= P(X_n = x_n \mid x_1 = x_1, \dots, X_{n-1} = x_{n-1}) P(X_1 = x_1, \dots, X_{n-1} = x_{n-1}) \\ &= \prod_{k=0}^{n-1} P(X_{k+1} = x_{k+1} \mid X_1 = x_1, \dots, X_k = x_k). \end{aligned}$$

In this specific kind of system where the probability of X_{k+1} only depends on the ones the step right before, we call it a *Markov Chain*, i.e. a probabilisitic version of dynamical systems.

§6 Day 6: Ulam's Problem; Chebyshev Inequality, Stirling Approximation, and Erdös-Renyi Random Graphs (Sep. 23, 2024)

Let us throw n balls into n boxes. Then consider N to be the number of empty boxes, and we have

$$P(N=k)=\frac{1}{n^n}\binom{n}{k}\sum_{\ell=0}^{k}(-1)^\ell\binom{k}{\ell}(k-\ell)^n,$$

with accompanying expectation

$$E[N]=\sum_{k=0}^{n}P(N=k).$$

However, we may simply the expression as follows; let N instead be written as a sum of indicators, i.e.

$$N=\sum_{i=1}^{N}1(\text{the } i\text{th box is empty}),$$

yielding

$$E[N]=\sum_{i=1}^{N}P(\text{the } i\text{th box is empty})=nP(\text{the } i\text{th box is empty})=n\frac{(n-1)^n}{n^n}.$$

Next example; the longest increasing subsequence of a random permutation. Let S_n be the set of bijections $\sigma:\{1,\dots,n\}\to\{1,\dots,n\}$. Then $|S_n|=n!$ where each σ has probability $\frac{1}{n!}$, and we define an increasing subsequence to be given by $\sigma(i_1)<\sigma(i_2)<\cdots<\sigma(i_k)$ for $i_1<i_2<\cdots<i_k$. Let $L_n(k)$ be the longest increasing subsequence of σ. Then L_n is a random variable; it remains to ask how big L_n is (Ulam's Problem). It is proven (though not in this class) that

$$\frac{L_n}{\sqrt{n}}\to 2$$

as $n\to\infty$, i.e. the expectation $E[L_n]\sim 2\sqrt{n}$ [9].

We prove a looser bound for now; let N_k be the number of increasing subsequences of length k. Then

$$N_k=\sum_{i_1<\cdots<i_k}1(\sigma(i_1)<\cdots<\sigma(i_k)),$$

and we may write

$$\begin{aligned}E[N_k]&=\sum_{i_1<\cdots<i_k}P(\sigma(i_1)<\cdots<\sigma(i_k))\\&=\binom{n}{k}\frac{1}{k!}=\frac{n!}{(n-k)!(k)!^2}.\end{aligned}$$

We now present Chebyshev's inequality,

$$P(X\geq x)\leq\frac{E[X1(X\geq x)]}{x},$$

where $x>0$. Using $X1(X\geq x)\geq x1(X\geq x)$, we have $E[X1(X\geq x)]\geq xP(X\geq x)$. Returning to earlier, we obtain $P(N_k>0)\leq E[N_k]=\frac{n!}{(n-k)!(k!)^2}$.

[9] for more, see here :3

Theorem 6.1 (Stirling's Formula)**.** $n! \sim \sqrt{2\pi} n^{n+\frac{1}{2}} e^{-n}$.

To prove this, start by considering the Gamma function $\Gamma(z) = \int_0^\infty t^{z-1} e^t \, dt$, which gives $\Gamma(n) = (n-1)!$. Let us consider the following;

$$\begin{aligned}
\Gamma(n+1) &= \int_0^\infty t^n e^{-t} \, dt \\
&= \int_0^\infty e^{n \log t - t} \, dt && \text{(Substitute } t = nx\text{)} \\
&= n e^{n \log n} \underbrace{\int_0^\infty e^{n(\log x - x)} \, dx}_{\text{of the form } \int e^{nf(x)} \, dx} \\
&\approx \int e^{nf(x^*) + \frac{f''(x^*)}{2}(x-x^*)^2} \, dx \\
&= e^{-n} \int e^{-\frac{n}{2}(x-1)^2} \, dx && \text{(Substitute } x = 1 + \tfrac{y}{\sqrt{n}}\text{)} \\
&= \frac{e^{-n}}{\sqrt{n}} \int_{-\infty}^\infty e^{-\frac{y^2}{2}} \, dy \\
&= \frac{e^{-n}}{\sqrt{n}} \sqrt{\int_{-\infty}^\infty \int_{-\infty}^\infty e^{-\frac{1}{2}(x^2+y^2)} \, dx \, dy} \\
&= \frac{e^{-n}}{\sqrt{n}} \sqrt{\int_0^\infty \int_0^{2\pi} e^{-\frac{r^2}{2}} r \, d\theta \, dr} \\
&= \frac{e^{-n}}{\sqrt{n}} \sqrt{2\pi \int_0^\infty r e^{-\frac{r^2}{2}} \, dr} \\
&= \frac{e^{-n}}{\sqrt{n}} \sqrt{2\pi}.
\end{aligned}$$

(I don't know where the hell this went. Oh well.)

We may now insert Stirling's formula into Ulam's problem to obtain

$$P_n(N_k > 0) \le \frac{n!}{(n-k)!(k!)^2} \approx c \frac{n^{n+\frac{1}{2}} e^{-n}}{k^{2k+1} e^{-2k} (n-k)^{n-k+\frac{1}{2}} e^{-(n-\frac{1}{2})}},$$

which cancels nicely. Using $k! \ge k^k e^{-k}$, we get $P(N_k > 0) \le (\frac{e\sqrt{n}}{k})^{2k}$. We may make the bound nicer by writing

$$\left(\frac{e\sqrt{n}}{k}\right)^{2k} \le \left(\frac{e}{3}\right)^{6\sqrt{n}}, \qquad (k = 3\sqrt{n})$$

and using $(\frac{e}{3})^6 \le e^{-\frac{1}{2}}$ yields that it is less than $e^{-\frac{n}{2}}$. Thus, we have that $P(L > \sqrt[3]{n}) \le \sum_{m=3\sqrt{n}}^\infty e^{-\frac{\sqrt{m}}{2}}$.

We now cover Erdös-Renyi random graphs. Let us have a graph on vertices $V = \{v_1, \ldots, v_n\}$, and edges $E = \{e_{ij}\} \subset V \times V$. We have that e_{ij} is in E with probability p, and not there with probability $1 - p$, considered independently over all undirected pairs (i, j).

Define a clique to be a complete subgraph of any graph $G(n, p)$, and let us have $\omega(G)$ to be the clique number, i.e. the size of the largest clique of G. This is approximately $C_p \log n$, where C_p is some constant. Let us have N_k as the number of cliques of size k. To calculate the expectation, let us have

$$N_k = \sum_{V' \subset V | V'| = k} 1_{\text{all } e_{ij} for all i,j \in V' \text{ exists}}.$$

Then we have

$$E[N_k] = \binom{n}{k} p^{\binom{k}{2}} =: f(k).$$

Observe that we have

$$\frac{f(k+1)}{f(k)} = \frac{\binom{n}{k+1} p^{\binom{k+1}{k}}}{\binom{n}{k} p^{\binom{k}{2}}} = \frac{n-k}{k+1} p^k.$$

Observing that $f(1) = n$ and $f(n) = p^{\frac{n(n-1)}{2}} << 1$, we see that f is unimodal. In particular, there is a unique point k_0 such that $f(k_0) \geq 1 > f(k_0 + 1)$. Thus,

$$\left(\frac{n}{k} - 1\right)^k p^{\frac{k(k-1)}{2}} \leq f(k) \leq n^k p^{\frac{k(k-1)}{2}}.$$

In particular, the right hand side is less than 1 if $np^{\frac{k-1}{2}} < 1$, and this evaluates out to $k > C_p \log n$. The left hand side is greater than 1 when $k \leq \frac{\log(\frac{n}{k} - 1)}{|\log p|} + 1$.

§10 Day 10: Erdös-Renyi Random Graphs and Cliques, Chebyshev Inequality, Moment Generating Function (Oct. 7, 2024)

We start at *Example 1.5.5* in Panchenko. Consider the Erdös-Renyi random graph, $G(n,p)$, where n is the number of vertices of the graph, and p is the probability that an edge is in the graph. Then a clique subset of $V=\{v_1,\dots,v_n\}$ is a complete graph, and we denote N_k to be the number of cliques of size k. We also define

$$f(k)=\mathbb{E}N_k=\binom{n}{k}p^{\binom{k}{2}},$$

i.e. $f(1)=n$, $f(n)=p^{\binom{n}{2}}$. Then k_0, defined as $f(k_0)\geq 1>f(k_0+1)$, has $k_0\sim c_p\log n$, where $c_p=\frac{2}{|\log p|}$. We also have

$$f(k+1)=\frac{n-k}{k+1}p^kf(k)\leq np^kf(k),\qquad f(k_0+m)\leq\frac{1}{n^m}(1-\varepsilon).$$

Then we also get bounds on k_0, which is greater than $\frac{(2-\varepsilon)\log n}{|\log p|}$, and so $p_0^k\leq\frac{1}{n^{2-\varepsilon}}$.

Using Chebyshev, we get that

$$\mathbb{P}(N_{k_0+m+1}>0)\leq\mathbb{E}[N_{k_0+m+1}]\leq\frac{1}{n^{m(1-\varepsilon)}}.$$

Observe that we have $\frac{(2-\varepsilon)\log n}{|\log p|}<k<k_0$, and $p^k\leq\frac{1}{n^{2-\varepsilon}}$, and

$$\frac{f(k+1)}{f(k)}=\frac{n-k}{k+1}p^k\leq np^k\leq\frac{1}{n^{1-\varepsilon}}.$$

We want to prove that there are lots of cliques there. To do this, we use the second moment method, $\mathbb{P}(|N_k-\mathbb{E}N_k|\geq x)\leq\frac{\operatorname{Var}(N_k)}{x^2}$. Using the fact that N_k is defined as

$$N_k=\sum_{\substack{W\subset V\\|W|=k}}1(W\text{ is a clique}),$$

we get that

$$\operatorname{Var}N_k=\sum_{W,W'}\operatorname{Cov}(1_W,1_{W'}).$$

If $|W\cap W'|\leq 1$, then $1_W,1_{W'}$ are clearly independent, so $\operatorname{Cov}(1_W,1_{W'})=0$. If W and W' share more than 1 vertex, though, call this number i. Then

$$\operatorname{Cov}(1_W,1_{W'})=\mathbb{E}[1'_W1_W]-\mathbb{E}[1_W]\mathbb{E}[1_{W'}]\leq\mathbb{E}[1'_W1_W]=\mathbb{P}(W,W'\text{ cliques}).$$

Note that there are $\binom{n}{k}\binom{k}{i}\binom{n-k}{k-i}$ pairs of W,W' with i vertices in common. Thus, we have that the probability of W,W' being cliques is given by

$$\mathbb{P}(W,W'\text{ cliques})=p^{\binom{k}{2}}p^{\binom{k}{2}-\binom{i}{2}}.$$

Then we may write

$$\begin{aligned}\operatorname{Var}(N_k)&\leq\sum_{i=2}^{k}\binom{n}{[}k]\binom{k}{i}\binom{n-k}{k-i}p^{\binom{k}{2}}p^{\binom{k}{2}-\binom{i}{2}}\\&=f^2(k)\sum_{i=2}^{k}\frac{\binom{k}{i}\binom{n-k}{k-i}}{\binom{n}{k}}p^{-\binom{i}{2}}.\end{aligned}$$

Using Chebyshev from earlier, we get

$$\mathbb{P}(|N_k - \mathbb{E}E_k| \geq \delta \mathbb{E}N_k) \leq \frac{1}{\delta^2} \sum_{i=2}^{k} a(i).$$

It remains to show that the right hand side is small for k in the range. This just leads to tedious calculation; to start, let us check $a(2)$:

$$\begin{aligned} a(2) &= \frac{\binom{k}{2}\binom{n-k}{k-2}}{\binom{k}{2}} p^{-\binom{2}{2}} \\ &= \frac{k^2(k-1)^2}{2} \frac{(n-k)\dots(n-2k+3)}{n(n-1)\dots(n-k+1)} \frac{1}{p} \\ &\leq \frac{k^4}{(n-k)^2} \frac{1}{2p} \leq c_p \frac{(\log n)^4}{n^2}. \end{aligned}$$

We also have that

$$a(k) = \frac{1}{\binom{n}{k}} p^{-\binom{k}{2}} = \frac{1}{f(k)}.$$

Let us define $b(i)$ as the ratio of two consecutive $a(i)$, and write

$$b(i) = \frac{a(i+1)}{a(i)} = \frac{\binom{k}{i+1}}{\binom{k}{i}} \frac{\binom{n-k}{k-i-1}}{\binom{n-k}{k-i}} p^{-\binom{i+1}{2}+\binom{i}{2}} = \frac{(k-i)^2}{(i+1)(n-2k+i+1)} p^{-i}.$$

We may derive the properties of a from b; in particular, we like b more because it is "nicer" than a. In particular, if $i \leq \frac{1}{3} \frac{\log n}{|\log p|}$, then $p^{-1} \leq n^{\frac{1}{3}}$, and we have $b(i) \sim n^{-\frac{2}{3}} \lessapprox 1$.

Likewise, if $i > \frac{3}{2} \frac{\log n}{|\log p|}$, then $p^{-i} > n^{\frac{3}{2}}$. Then $b(i) \gtrapprox \frac{n^{\frac{1}{2}}}{k} > 1$. We have that $\frac{b(i+1)}{b(i)} \gtrapprox \frac{1}{p} > 1$, so b is increasing.

We now return to sums of independent random variables. Let $X_1, \dots, X_n$ be i.i.d., and $\mathbb{E}X_i^2 < \infty$. Let us have $\overline{X_n} = \frac{1}{n} \sum_{i=1}^{n} x_i$, and

$$\mathbb{P}(\left|\overline{X_n} - EEX_1\right| > \varepsilon) = \frac{\operatorname{Var} \overline{X_n}}{\varepsilon^2} = \frac{\sigma^2}{n\varepsilon^2},$$

where we let ε^2 be the variance of X_i. We could do better, i.e.

$$\mathbb{P}(X \geq t) \leq \frac{\mathbb{E}[e^{\lambda x}]}{e^{\lambda t}}.$$

This gives

$$\mathbb{P}(X_1 + \dots + X_n \geq t) \leq \frac{\mathbb{E}[e^{\lambda(X_1+\dots+X_n)}]}{e^{\lambda t}} = \frac{EE[e^{\lambda X_1}]^n}{e^{\lambda t}}.$$

This is called exponential Chebyshev. For example, let $X = \sum \varepsilon_i a_i$, where ε_i are independent Rademacher with $\mathbb{P}(\varepsilon_i = 1) = \frac{1}{2} = \mathbb{P}(\varepsilon_i = -1)$ (read: coinflip distribution). Then we may apply exponential Chebyshev and optimize to get (Theorem 3.1 in Panchenko)

$$\mathbb{P}\left(\sum_{i=1}^{n} \varepsilon_i a_i \geq t\right) \leq e^{-\frac{t^2}{2\sum_{i=1}^{n} a_i^2}}.$$

We also have

$$\mathbb{P}\left(\left|\sum_{i=1}^{n} \varepsilon_i a_i\right| \geq t\right) \leq 2e^{-\frac{t^2}{2\sum_{i=1}^{n} a_i^2}}.$$

By the law of large numbers for a fair coin, we may take $a_i = \frac{1}{n}$ as per $\overline{X_n}$, and we get

$$\mathbb{P}\left(\left|\frac{1}{n}\sum_{i=1}^{n}\varepsilon_i\right| \geq t\right) \leq 2e^{-\frac{nt^2}{2}}. \quad \square$$

Small digression; let $X_1, \dots$ be i.i.d., and consider $\mathbb{E}[e^{\lambda x}] < \infty$. We call the LHS $M(\lambda)$, i.e. the moment generating function. In particular, $M'(0) = \mathbb{E}[X]$. We also have $M''(0) = \mathbb{E}[X^2] = \sum_{n=0}^{\infty} \frac{\lambda^n}{n!}\mathbb{E}[X^n]$, meaning that we have $X^n \leq C_\lambda[e^{\lambda X} + e^{-\lambda X}]$. Notice that we may write

$$\mathbb{P}\left(\frac{X_1 + \cdots + X_n}{n} > t\right) \leq \frac{M(\lambda)^n}{e^{\lambda t n}} = e^{-n\{\lambda t - \log M(\lambda)\}}.$$

Optimizing this over λ, we get that the LHS is less than or equal to $e^{-nI(d)}$, where $I(t) = \sup_\lambda\{\lambda t - \log M(\lambda)\}$.

Let $a_1^2 + \cdots + a_n^2 = 1$. Then $\mathbb{E}X^{2k+1} = 0$, and we may write

$$\begin{aligned}
\mathbb{E}X^{2k} &= \int_0^\infty 2kt^{2k-1}\mathbb{P}(|X| > t)\,dt \\
&\leq \int_0^\infty 2kt^{2k-1}2e^{-\frac{t^2}{2}}\,dt \\
&= k2^{k+1}\int_0^\infty u^{k-1}e^{-i}\,du \\
&= k2^{k+1}\Gamma(k) \\
&= k2^{k+1}(k-1)! = 2^{k+1}k!.
\end{aligned}$$

Let $Z \sim N(0,1)$. Then we have

$$\mathbb{E}[e^{\lambda z^2}] = \int_{-\infty}^{\infty} e^{\lambda z^2 - \frac{z^2}{2}}\,\frac{dz}{\sqrt{2\pi}}. < \infty$$

if and only if $\lambda < \frac{1}{2}$. Let $Y = X^2 - 1$. Then $e^x \leq 1 + x + \frac{x^2}{2} + \sum_{k=3}^{\infty}\frac{x_+^k}{k!}$. We denote $x_+ = x$ if $x \geq 0$, and 0 otherwise. $\square$

§12 Day 12: Inequalities (Oct. 16, 2024)

Let X_i be independently distributed B(p), and consider $X_n = \frac{1}{n}\sum_{i=1}^n X_i$. Then (where $q > p$)

$$\mathbb{P}\left(\overline{X_n} \geq q\right) \leq e^{-n\lambda q}\mathbb{E}[e^{\lambda x_i}]^n = e^{-n\left\{\lambda q - \log(1-p\pm pe^\lambda)\right\}}.$$

We want to take the supremum of $\{\lambda q - \log(1-p+pe^\lambda)\}$; then we have

$$f'(\lambda) = q - \frac{pe^\lambda}{1-p+pe^\lambda},$$

where we note if $f'(\lambda) = 0$, then $q(1-p) + qpe^\lambda = pe^\lambda$ implies $e^\lambda = \frac{q(1-p)}{p(1-q)} > 1$. We write $D(q\,||\,p)$ to be the Kullback-Leibler divergence, aka relative entropy, i.e. $H(q\,||\,p)$. Let M be a differentiable manifold of dimension n. Then let $D : M \times M \to [0,\infty)$, and consider M as a parameterized family of probability measures. Then

$$\begin{aligned} D(q,p) &\geq 0; \\ D(q,p) &= 0; && \text{(if and only if } q = p) \\ D(p, p+dp). && \text{(should be positive def. quadratic in } dp) \end{aligned}$$

We check that this is indeed true.

(a) We check that D is non-negative.

$$\begin{aligned} &q\log\frac{q}{p} + (1-q)\log\frac{1-q}{1-p} \\ &= q\left(-\log\frac{p}{q}\right) + (1-q)\left(-\log\frac{1-p}{1-q}\right) \\ &\geq q\left(1-\frac{p}{q}\right) + (1-q)\left(1-\frac{1-p}{1-q}\right) \\ &= q - p + 1 - q - (1-p) = 0. \end{aligned}$$

(b) We now check that it is identically zero iff $q = p$. Intuitively, there is no entropy needed to move $p \to q$ if they are equal.

$$\begin{aligned} \inf_q \underbrace{\left\{q\log\frac{q}{p} + (1-q)\log\frac{1-q}{1-p}\right\}}_{g'(a)} \\ g'(a) = \log\frac{q}{p} - q\frac{p}{q}\frac{1}{p} - \log\frac{1-q}{1-p} = 1 \\ \implies \log\frac{q}{p} = \log\frac{1-q}{1-p}, \end{aligned}$$

which occurs only when $p = q$.

(c) We leave the third alone for now.

An application of this is to classification algorithms. Consider t a classifier, and E_n an empirical error, i.e.

$$E_n(t) = \frac{1}{n}\sum_{i=1}^m L(y_i, f(x_i)),$$

where L is some loss function, and we consider (X_i, Y_i) i.i.d. $\mathcal{F} = \{f_1, \ldots, f_N\}$. The generalization error $E(f) = E[L(X, f(y))]$, and suppose we have

$$\mathbb{P}(E_n(f) \geq E(f) + \varepsilon) < e^{-cn\varepsilon^2}.$$

Then $\mathbb{P}(\forall f \in \mathcal{F}, E(f) \leq E_n(f) + \varepsilon) \geq 1 - Ne^{cn\varepsilon^2}$. Let $\delta = Ne^{-Cn\varepsilon^2}$; to get confidence $1 - \delta$ that the generalization error is within ε; then we need $n = \frac{1}{c\varepsilon^2} \log \frac{N}{\delta}$.

§17 Day 17: Distributions Related to Gaussian (Nov. 11, 2024)

This is *Section 4.4* in Panchenko! Recall that the Γ function is given by

$$\Gamma(\alpha) = \int_0^\infty x^{\alpha-1}e^{-x}\,dx;$$

if we divide both sides by $\Gamma(\alpha)$ and perform a change of variables $x = \beta y$ for $\beta \geq 0$, we get

$$1 = \int_0^\infty \frac{1}{\Gamma(\alpha)}x^{\alpha-1}e^{-x}\,dx = \int_0^\infty \frac{\beta^\alpha}{\Gamma(\alpha)}y^{\alpha-1}e^{-\beta y}\,dy;$$

thus, we see that

$$f_{\alpha,\beta}(x) = \frac{\beta^\alpha}{\Gamma(\alpha)}x^{\alpha-1}e^{-\beta x}1_{\{x\geq 0\}}$$

for each $\alpha, \beta > 0$ is a density, and is called the *Gamma distribution with parameters* α, β, and is written $\Gamma(\alpha, \beta)$. Let $X_i \sim \Gamma(\alpha_i, \beta)$ be independent; then we have $X_1 + \cdots + X_n \sim \Gamma(\alpha_1 + \cdots + \alpha_n, \beta)$. If X, Y are independent with densities f, g, then $X + Y$ has density

$$(f * g)(x) = \int_{-\infty}^\infty f(x-y)g(y)\,dy.$$

We use these two properties to inductively prove that the sum is indeed distributed $\Gamma(\alpha_1 + \cdots + \alpha_n, \beta)$; for $n = 2$, we have (from convolution directly),

$$\begin{aligned}
&\int_0^x \frac{\beta^{\alpha_1+\alpha_2}}{\Gamma(\alpha_1)\Gamma(\alpha_2)}e^{-\beta(x-y)}e^{-\beta y}(x-y)^{\alpha_1-1}y^{\alpha_2-1}\,dy \\
&= \frac{\beta^{\alpha_1+\alpha_2}}{\Gamma(\alpha_1)\Gamma(\alpha_2)}e^{-\beta x}\int_0^x (x-y)^{\alpha_1-1}y^{\alpha_2-1}\,dy && \text{(Substitute } y = xz\text{)} \\
&= \frac{\beta^{\alpha_1+\alpha_2}}{\Gamma(\alpha_1)\Gamma(\alpha_2)}e^{-\beta x}\int_0^1 (x-xz)^{\alpha_1-1}xz^{\alpha_2-1}x\,dz \\
&= \frac{\beta^{\alpha_1+\alpha_2}}{\Gamma(\alpha_1)\Gamma(\alpha_2)}e^{-\beta x}x^{\alpha_1+\alpha_2-1}\int_0^1 (1-z)^{\alpha_1-1}z^{\alpha_2-1}\,dz \\
&= \beta^{\alpha_1+\alpha_2}e^{-\beta x}x^{\alpha_1+\alpha_2-1}C1_{\{x\geq 0\}},
\end{aligned}$$

where

$$C = \frac{\int_0^1 (1-z)^{\alpha_1-1}z^{\alpha_2-1}\,dz}{\Gamma(\alpha_1)\Gamma(\alpha_2)} = \frac{1}{\Gamma(\alpha_1+\alpha_2)}.$$

We see that this just means that by induction, we get a $\Gamma(\alpha_1 + \cdots + \alpha_n, \beta)$ distribution as desired. Now, if $g_1, \ldots, g_n$ are independent standard Gaussians, then $g_1^2 + \cdots + g_n^2 \sim \chi_n^2$, i.e. "chi squared with n degrees of freedom", where

$$\chi_1^2 \sim \frac{1}{\sqrt{2n}}x^{\frac{1}{2}-1}e^{-\frac{1}{2}}1_{\{x\geq 0\}} \sim \Gamma\left(\frac{1}{2}, \frac{1}{2}\right).$$

We know that $\chi_n^2 \sim \Gamma(\frac{n}{2}, \frac{1}{2})$; if $X \sim \chi_k^2$, $Y \sim \chi_m^2$, then the distribution of the ratio is given by $Z = \frac{X/k}{Y/m} \sim F_{k,m}$, i.e. the "$F$ distribution with degrees of freedom k, m".

Lemma 17.1. If $X, Y > 0$ and independent with densities f, g then $\frac{X}{Y}$ has density $\int_0^\infty f(xy)g(y)y\,dy$.

Write

$$\begin{aligned}\mathbb{P}\left(\frac{X}{Y}\le t\right)=\mathbb{P}(X\le tY)&=\int_0^\infty \mathbb{P}(X\le ty)g(y)\,dy\\&=\int_0^\infty\int_0^{ty} f(x)g(y)\,dx\,dy && \text{(Substitute } x=zy)\\&=\int_0^\infty\int_0^{t} f(zy)g(y)y\,dz\,dy && \text{(Fubini)}\\&=\int_0^t\left(\int_0^\infty f(zy)g(y)y\,dy\right)dz\end{aligned}$$

as desired. Now, write $X\sim\chi^2_n$, $Y\sim\chi^2_m$, and $f_{\chi^2_k}(xy)$. We have

$$\begin{aligned}f_{\frac{X}{Y}}(x)&=\int_0^\infty \underbrace{\frac{\left(\frac12\right)^{\frac k2}}{\Gamma(\frac k2)}(xy)^{\frac k2-1}e^{-\frac12 xy}}_{f_{\chi^2_k(xy)}}\underbrace{\frac{\left(\frac12\right)^{\frac m2}}{\Gamma(\frac m2)}(xy)^{\frac m2-1}e^{-\frac12 y}}_{y_{\chi^2_m}(y)}y\,dy\\&=\frac{\left(\frac12\right)^{\frac{k+m}{2}}}{\Gamma(\frac k2)\Gamma(\frac m2)}x^{\frac k2-1}\int_0^\infty y^{\frac{x+k}{2}-1}e^{-\frac12(x+1)y}\,dy && (z=\tfrac12(x+1)y)\\&=\frac{\Gamma(\frac{k+m}{k})}{\Gamma\frac k2\Gamma(\frac m2)}x^{\frac k2-1}(1+x)^{-\frac{k+m}{2}}.\end{aligned}$$

In particular, this means

$$f_{k,m}(x)=\frac{\Gamma(\frac{k+m}{2})}{\Gamma(\frac k2)\Gamma(\frac k2)}k^{\frac k2}m^{\frac m2}x^{\frac k2-1}(m+kx)^{-\frac{k+m}{2}}.$$

Now, let $g_0,\dots,g_n$ be independent standard Gaussians. Then the distribution of

$$T=\frac{g_0}{\sqrt{\frac1n(g_1^2+\cdots+g_n^2)}}$$

is the Student's T-distribution with n degrees of freedom, often written t_n. Writing

$$T^2=\frac{g_0^2}{\frac1n(g_1^2+\cdots+g_n^2)}\sim F_{1,n},$$

we have that $\mathbb{P}(T^2\le t^2)=\mathbb{P}(-t<T<t)=2\mathbb{P}(0\le T\le t)$ by symmetry; we may write

$$2\int_0^t f_T(x)\,dx=\int_0^{t^2} f_{1,n}(x)\,dx=\int_0^t f_{1,n}(y^2)2y\,dy,$$

so

$$f_T(t)=f_{1,n}(t^2)t=\frac{\Gamma(\frac{n+1}{2})}{\Gamma(\frac12)\Gamma(\frac n2)}\frac{1}{\sqrt n}\left(1+\frac{t^2}{n}\right)^{-\frac{n+1}{2}}.$$

As $n\to\infty$, we get that

$$\left(1+\frac{t^2}{2\left(\frac n2\right)}\right)^{-\frac n2+\frac12}\xrightarrow{n\to\infty}e^{-\frac{t^2}{2}},$$

and

$$\frac{\Gamma(\frac{n+1}{2})}{\Gamma(\frac n2)}\frac{1}{\sqrt n}\xrightarrow{n\to\infty}\frac{1}{\sqrt{2\pi}},$$

so we have

$$\lim_{n\to\infty}f_T(t)=\frac{1}{\sqrt{2\pi}}e^{-t^2/2}.$$

We now move onto linear regressions (section 4.5 in Panchenko). Let $(x_1, y_1), \ldots, (x_n, y_n)$ be data points; the simple linear regression (SLR) model is $Y_i = \beta_0 + \beta_1 X_i^2 + \varepsilon_i$, where X_i are independent variables, $y = f(x) = \beta_0 + \beta_1 x$ is the regression line, and ε_i are Gaussian distributed $N(0, \sigma^2)$. The density of $\vec{y}$ is given by

$$\ell_{\beta_0,\beta_1,\sigma} = \prod_{i=1}^{n} \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2\sigma^2}(y-\beta_0-\beta_i X_i)^2},$$

where ℓ represents the *likelihood function*. We may write

$$\max_{\beta_0,\beta_1,\sigma} \ell_{(\beta_0,\beta_1,0,(x_1,\ldots,x_n,y_1,\ldots,y_n))}$$

as the maximum likelihood estimate. We start by maximizing over β_0, β_1. Now, we just need to minimize $L = \sum_{i=1}^{n}(y_i - \beta_0 - \beta_i x_i)^2$. We have

$$\frac{\partial L}{\partial \beta_0} = -2\sum_{i=1}^{n} Y_i - (\beta_0 + \beta_1 X_i) = 0, \frac{\partial L}{\partial \beta_1} \qquad = -2\sum_{i=1}^{n} Y_i - (\beta_0 + \beta_1 X_i)X_i = 0.$$

Solving the above, we have

$$\hat{\beta}_0 := \overline{Y} - \hat{\beta}_1 \overline{X}, \qquad \hat{\beta}_1 := \frac{\overline{XY} - \overline{XY}}{\overline{X^2} - \overline{X}^2},$$

where $\overline{X}, \overline{Y}, \overline{XY}$ are given by $\frac{1}{n}\sum_{i=1}^{n} X_i$, $\frac{1}{n}\sum_{i=1}^{n} Y_i$, and $\frac{1}{n}\sum_{i=1}^{n} X_i Y_i$ respectively. Now, we want to maximize

$$-n\log\sigma - \frac{1}{2\sigma^2}\sum_{i=1}^{n}(y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2$$

over σ. A few more things that I just didn't record cuz too tired. x3

§21 Day 21: Markov Chains (Nov. 25, 2024)

Assume that our state space is finite. The period d_i of a state s_i is defined by

$$d_i = \gcd\{n \geq 1 \mid p_{ii}(n) > 0\}.$$

If $d_i = 1$, then s_i is called *aperiodic*. A Markov chain is called *irreducible* if each pair of states communicates.

Lemma 21.1. If a Markov chain is irreducible and aperiodic, then there exists an N so that, for all $n \geq N$, and i, j, $p_{ij}(n) > 0$.

Let d be the smallest integer greater than 1 such that d can be written $a_1 n_1 + \cdots + a_k n_k$, where a_i's are integers (not necessarily positive), and n_i's are all in $T(s_j) = \{n \geq 1 \mid p_{11}(n) > 0\}$, i.e. the set of all times the chain starting at s_1 can return to s_1. Start by claiming that $d = 1$. If d divides all $n \in T(s_1)$, then $d = 1$ clearly. Supposing not, then $n = ad + r$ with $1 \leq r < d$. But $r = n - ad = n - a(a_1 n_1 + \cdots + a_k n_k)$, contradicting the minimality of d. So we learn that $1 = a_1 n_1 + \cdots + a_k n_k$.

Taking the largest $|a_k|$, which we may as well call $|a_1|$, we have that

$$N_1 = |a_1|\, n_1(n_1 + \cdots + n_k).$$

We claim that for any $n \geq N_1$, $n = c_1 n_1 + \cdots + c_k n_k$, where $c_k \geq 0$ are integers. Let $n = N_1 + \ell$, with $1 \leq \ell < n_1$. Then

$$n = N_1 + \ell = |a_1|\, n_1(n_1 + \cdots + n_k) + \ell(a_1 n_1 + \cdots + a_k n_k) = c_1 n_1 + \cdots + c_k n_k.$$

In particular, $c_1 = |a_1|\, n_1 + \ell a_i \geq \ell(|a_1 + a_i|) \geq 0$. Letting $N_1 + n_1 = (|a_1|\, n_1 + 1)n_1 + |a_1|\, n_1(n_2 + \cdots + n_k)$, $n = N_1 + n_1 + \ell$. Thus, we have $p_{11}(n) \geq (p_{11}(n_1))^{c_1} \ldots (p_{11}(n_k))^{c}_{k} > 0$. So for all $n \geq N_1, p_{11}(n) > 0$, and so on for all N_m and $p_{mm}(n) > 0$ respectively. Let $N' = \max\{N_1, \ldots, N_m\}$. Then $p_{ii}(n) > 0$ if $n \geq N'$. For all i, j, there exists k such that $1 \leq k \leq m$, with $p_{ij}(k) > 0$. Thus, $N = N' + m$, then for all $n \geq N$, $p_{ij}(n) > 0$. □

Last week, we showed that there is always at least one stationary distribution. If the state space S were not finite, then we may take $S = \mathbb{Z}$ and $p_{i,i+1} = 1, p_{ii} = 0$ to see that there is no stationary distribution.

Lemma 21.2. If a Markov chain is irreducible, then the statinoary distribution is unique.

To start, with, we have that $(P - I)(1) = 0$. We now show that irreducible implies that $\operatorname{rank}(P - I) = m - 1$. Let v_z be the largest index in v; then

$$v_z = \sum_{j=1}^{m} p_{zj} v_j \leq \sum_{j=1}^{m} p_{zj} v_z.$$

In particular, $p_{zj} v_j = p_{zj} v_z$ for all j, and all v_j's with $p_{zj}(z) > 0$ are also equal to p_z. By irreducibility, there exists some k such that $p_{zj}(k) > 0$; then $\ker(P - I)^t$ is one dimensional if and only if the stationary distribution is unique. We may then write

$$\frac{1}{n}\left(1 + P + P^2 + \cdots + P^{n-1}\right) \to A = \begin{pmatrix} \vec{\mu} \\ \vdots \\ \vec{\mu} \end{pmatrix},$$

so $\lim_{n\to\infty}$ of the above gives the desired A. □

Let $f : S \to \mathbb{R}$;

$$\lim_{n\to\infty} \mathbb{E}\left[\frac{1}{n}\sum_{k=1}^{n} f(X_k)\right] = \sum_{i=1}^{m} f(i)\mu_i.$$

Then $f = \sum f(X_i)1_S$. By the law of large numbers,

$$\lim_{n\to\infty} \frac{1}{n}\sum_{k=0}^{n-1} f(X_k) = \sum_{i=1}^{m} f(s_i)\mu_i.$$

This is called the ergodic theorem.

Theorem 21.3. If a Markov chain is aperiodic and irreducible, then

$$\lim_{n\to\infty} p^n = \begin{pmatrix} \vec{\mu} \\ \vdots \\ \vec{\mu} \end{pmatrix}$$

if $v_i = P(X_0 = S_1)$, then $\lim_{n\to\infty} P(X_n = s_j)) = \mu_j$.

P contracts L^1 norm on $\mathbb{R}^m$; $\|x\|_1 = \sum_{i=1}^m |x_i|$. Then

$$\|xP - yP\|_1 = \sum_{i=1}^{m}\left|\sum_{i=1}^{m} p_{ij}(x_i - y_i)\right| \le \sum_{j=1}^{m}\sum_{i=1}^{m} p_{ij}\,|x_i - y_i| = \sum_{i=1}^{m} |x - i - y_i| = \|x - y\|_1.$$

Then there exists N such that $p_{ij}(N) > 0$; for all i, j, let ε be the minimum over i, j, so $p_{ij}(N) \ge \varepsilon > 0$. Then

$$\begin{aligned}\|xP^n - yP^n\|_1 &= \sum_{j=1}^{m}\left|\sum_{i=1}^{m}(pij(N) - \varepsilon)(x_i - y_i)\right| \\ &\le \sum_{j=1}^{m}\sum_{i=1}^{m}(p_{ij}(N) - \varepsilon)\,|x_i - y_i| = (1 - m\varepsilon)\,\|x - y_1\|.\end{aligned}$$

Thus, we have that $\left\|vP^N - \mu\right\|_1 \le (1 - m\varepsilon)\,\|v - \mu\|_1$. In particular, $\left\|vP^{Nk} - \mu\right\|_1 \le (1 - m\varepsilon)^k\,\|v - \mu\|_1$; then

$$\|vP^n - \mu\|_1 \le (1 - m\varepsilon)^{\left\lfloor \frac{n}{N} \right\rfloor}\,\|v - \mu\|_1 \to 0$$

as $n \to \infty$. □

We say that M is reversible if

$$\mu_i p_{ij} = \mu_j p_{ji}.$$

This is a stronger condition than being stationary. We check this by proving that if M is reversible, then it is stationary;

$$\sum_i \mu_i p_{ij} = \sum_i \mu_j p_{ji} = \mu_j$$

then

$$\begin{aligned} P(X_0 = s_i, X_1 = s_j) &= \mu_i p_{ij}, \\ P(X_0 = s_j, X_1 = s_i) &= \mu_j p_{ji}, \\ P(X_0 = s_i, X_1 = s_j, X_2 = s_k) &= \mu_i p_{ij} p_{jk}, \\ P(X_0 = s_k, X_1 = s_j, X_2 = s_i) &= \mu_k p_{kj} p_{ji}, \end{aligned}$$

and so on. For an example, let $S = \{1, \dots, N\}$, with the equivalence relation of modulo N, with $p_{i,i+1} = p$, $p_{i,i-1} = 1-p$, and $p_{ij} = 0$ otherwise. Then the stationary distribution (invariant measure) is $\mu_{i-1}p_{i-1,1} + \mu_{i+1}p_{i+1,i} = \mu_i$, and we may pick $\mu_i = \frac{1}{N}$.

Another example; let V be a set of vertices (read: states), and $E \subset V \times V$. Let n_i be the number of neighbors of s_i, and let

$$p_{ij} = \begin{cases} \frac{1}{N_i} & \text{if } (s_i, s_j) \in E, \\ 0 & \text{otherwise.} \end{cases}$$

Irreducibility in graphs is equivalent to the graph connectivity condition, i.e. you can get from anywhere to anywhere else in a finite number of steps. *mu* is reversible means that $\mu_i \frac{1}{N_i} = \mu_j \frac{1}{N_j}$, which is equal to some constant independent of i, j. Said constant C is given by $\sum \mu_i = C \sum N_i$, so the constant is just the reciprocal of the sum of neighbors; so

$$\mu_i = \frac{N_i}{N}$$

is reversible (and is thus stationary). $\square$

§23 Day 23: Markov Chains, Pt. 2 (Dec. 2, 2024)

Let $S = \{s_1, \dots, s_m\}$ be a finite set of states; consider the chain $X_0 = s_1, X_1, \dots$, and let $T = \min\{n \geq 1 \mid X_n = S_1\}$ be the "return time". If the Markov chain does not start at s_1, this is called a "hitting time" of s_1. In particular, it is an example of a *stopping time*, i.e. T is a stopping time if $1\{T = n\} = f_n(X_0, \dots, X_n)$. Note that

$$\begin{aligned} S &= \text{last time we visited} s_1 \text{ is not a stopping time.} \\ U &= T - 1 \text{ is not a stopping time.} \end{aligned}$$

We may write $P(X_{k+1} = x_{k+1} \mid X_0 = x_0, \dots, X_k = x_k) = P(X_{k+1} = x_{k+1} \mid X_k = x_k)$. The Markov property generalizes to functions (??). If T is a stopping time with $T < \infty$, then $P(y(X_T + \dots) \mid X_0 = x_0, \dots, X_T = x_T) = P(g(X_0 + \dots) \mid X_0 = x_T)$. We now prove it for hitting times.

$$\begin{aligned} &P(g(X_T + \dots) = a \mid X_0 = x_0, \dots, X_T = s_1) \\ &= P(g(X_0 + \dots) = a \mid X_0 = s_1). \end{aligned}$$

This is called the *strong Markov property*. Directly write as follows,

$$\begin{aligned} &P(g(X_T + \dots) = a \mid X_0 = x_1, \dots, X_T = s_1, T = n) \\ &= P(g(X_n + \dots) = a, X_0 = x_1, \dots, X_n = S_1, T = n) \\ &= P(g(X_n + \dots) = a \mid X_n = s_1) P(X_0 = x_1, \dots, X_n = s_1, T = n). \end{aligned}$$

Summing over n, we get that $P(g(X_T + \dots) = a, X_0 = x_1, \dots, X_T = s_1) = P(g(X_0 + \dots) = a \mid X_0 = s_1) P(X_0 = x_1, \dots, X_T = s_1)$. Recall that we showed that if a Markov chain is aperiodic and irreducible, then

$$\lim_{n \to \infty} \mathbb{E}\left[\frac{1}{n} \sum_{k=1}^{m} 1(X_k = s_i)\right] = \mu_j$$

in time t up to n it returns to s_1, at times $T_1, T_1 + T_2, T_1 + T_2 + T_3, \dots, T_1 + \dots + T_L$. L is the number of times that you came back to s_1 in $\{1, \dots, n\}$. Note that $T_1, \dots, T_L$ are independent and identically distributed. Consider $N_1(j), \dots, N_L(j)$, where $N_i(j)$ are the number of vists to s_j on $T_1 + \dots + T_i$, $T_1 + \dots + T_{i+1}$. Also note that each $N_i(j)$ is i.i.d., except we just don't know the distributino. Then

$$N_1(j) + \dots + N_L(j) \sim \sum_{i=1}^{n} 1(X_i = s_j); \qquad \frac{1}{n} \sum_{k=1}^{n} 1(X_k = s_j) \sim \frac{N_1(j) + \dots + N_L(j)}{T_1 + \dots + T_L}.$$

As $n \to \infty$, we have that

$$\frac{\mathbb{E}N_i(j)}{\mathbb{E}T_i} = \mu_j,$$

which implies that

(i) $\lim_{n \to \infty} \sum_{k=1}^{n} f(X_k) = \sum_j f(s_j)\mu_j$. This is an Ergodic theorem, and this is true for any f on S.

(ii) $\mu_j = \dfrac{\mathbb{E}N_1(j)}{\mathbb{E}T_1}$. In particular, any state works.

We now introduce the Metropolis algorithm. Let $S = \{s_1, \dots, s_m\}$ be a graph, where s_i has N_i neighbors (which we say $s_i \sim s_j$ if they are neighbors). Let

$$p_{ij} = \begin{cases} \frac{1}{N_i} \min\{\frac{\mu_j N_i}{\mu_i N_j}, 1\} & s_j \sim s_1, \\ 1 - \sum_{s_j \sim s_i} p_{ij} & i = j, \\ 0 & \text{otherwise.} \end{cases}$$

This chain has μ_i as its stationary (reversible) distribution, i.e.

$$\frac{1}{n} \sum_{k=1}^{\infty} f(X_k) \to \sum_i f(s_i)\mu_i,$$
$$\frac{1}{n} \sum_{k=1}^{\infty} 1(X_k = s_i) \to \mu_i.$$

I'm not recording the Ising mnodel. Not sure where we're going with that.

Consider the riffle shuffle; cut a deck of n cards into $c, n-c$ with probability $\binom{n}{c}2^{-n}$. If A cards are on the left, B are on the right, the chance of the next card from the left is $\frac{A}{A+B}$, and $\frac{B}{A+B}$.